# Sai Sankar Bhuvanapalli

850-354-0461 | sankar156b@gmail.com | linkedin.com/in/saisankarb | github.com/saisankar20 | leetcode.com/u/sb23w

## Education

**Florida State University** — Tallahassee, FL
*Masters of Science in Computer Science GPA: 3.85/4.00* — *Aug 2023 – May 2025*
- Relevant Coursework: Algorithms, Advanced Operating Systems, Machine Learning, Advanced Data Science, Analytic Methods

**Jawaharlal Nehru Technological University** — Hyderabad, India
*Bachelor of Technology in Computer Science & Engineering GPA: 3.13/4.00* — *Aug 2017 – Jul 2021*

## Skills

**Languages:** Python, Java, SQL, Bash, react
**Backend & APIs:** FastAPI, Spring Boot, REST APIs
**Data Engineering:** PySpark, Databricks, Snowflake, SSIS
**AI/ML & RAG:** LangChain, Pinecone, SentenceTransformers, Gemini API
**Streaming & Messaging:** Apache Kafka, AWS Kinesis, Amazon SQS, Apache Airflow
**Databases & Warehouses:** Amazon Redshift, PostgreSQL, MySQL, DynamoDB, Hive, HBase
**Cloud/Infra:** AWS (Lambda, S3, Glue, EMR, DynamoDB, Kinesis), Kubernetes, Docker, Terraform, CloudFormation
**Observability & BI:** Amazon CloudWatch, Prometheus, Grafana, Power BI, QuickSight, Tableau
**DevOps & Practices:** Jenkins, GitHub Actions, Git, Linux, CI/CD, Agile/Scrum, unit testing, data modeling

## Experience

**Florida Power & Light** — Hyderabad, India
*Data Engineer* — *Jul 2021 – Aug 2023*
- Designed and maintained high-reliability RESTful microservices in Java/Spring Boot, achieving 99% uptime and supporting 1M+ monthly API calls across distributed systems.
- Implemented event-driven data services using AWS Lambda, DynamoDB, and S3, improving latency by 25% and reducing downtime incidents by 40%.
- Led migration of legacy FTP data pipelines to secure SFTP with end-to-end encryption, ensuring integrity across multi-service architectures. Collaborated with cross-functional teams to define API contracts, optimize query performance, and drive observability and fault-tolerance improvements
- Built CI/CD pipelines and containerized deployments with Jenkins, Docker, and CloudFormation, enabling seamless releases and zero-downtime deployments.

**Florida State University – Information Technology Services** — Remote
*Data Intern* — *May 2024 – Dec 2024*
- Developed serverless data APIs using Python and AWS Lambda, enabling real-time ingestion and transformation across distributed microservices.
- Built ETL services and API endpoints feeding live dashboards for 5+ departments, improving data accessibility and latency.
- Integrated AIM system APIs into AWS Redshift, streamlining backend data models and reducing manual reporting by 60%.
- Implemented error-tolerant data validation services with automated recovery logic to enhance system reliability and audit transparency.

**Saayam for All** — Remote
*Full-Stack Developer* — *Jul 2025 – Present*
- Architected FastAPI-based backend services containerized with Docker and Kubernetes, creating volunteer-matching APIs..
- Designed GraphQL-compatible REST endpoints for geospatial search using Uber H3 and DynamoDB, enabling sub-second location lookups.
- Automated infrastructure provisioning via CloudFormation, ensuring reproducibility and secure rollbacks across environments.

## Projects

**Portfolio Risk Analytics Dashboard** | GitHub
- Engineered backend modules in Python for financial analytics and data aggregation, exposing APIs for real-time visualization.
- Implemented async computation services and modular architecture to support scalable dashboard queries.
- Ensured reliability through a comprehensive pytest suite and continuous integration with GitHub Actions.

**RAG Roadmap Generator** | Website | GitHub
- Developed serverless GraphQL-style APIs with Netlify Functions + Python backend, generating learning roadmaps.
- Deployed asynchronous RAG pipelines using SentenceTransformers and LangGraph, optimizing compute efficiency.
- Built an interactive React-based frontend for dynamic roadmap visualization, integrating API responses into responsive UI components for real-time user interaction.

**Reddit Trends Analysis Pipeline** | GitHub
- Ingested 10M+ Reddit posts/day into PostgreSQL and Redshift with automated cleaning and normalization in Python.
- Scheduled Apache Airflow DAGs to manage dependencies, reducing manual effort by 90% and cutting query latency by 40%.
- Delivered durable pipelines and warehouse models to support downstream analytics and reporting.